

Due October 25, 6:00pm

1. (20 pts) A greedy algorithm—so to speak

The founder of LinkedIn, the professional networking site, decides to crawl LinkedIn's relationship graph to find all of the *super-schmoozers*. (He figures he can make more money from advertisers by charging a premium for ads displayed to super-schmoozers.) A *super-schmoozers* is a person on LinkedIn who has a link to at least 20 other super-schmoozers on LinkedIn.

We can formalize this as a graph problem. Let the undirected graph $G = (V, E)$ denote LinkedIn's relationship graph, where each vertex represents a person who has an account on LinkedIn. There is an edge $\{u, v\} \in E$ if u and v have listed a professional relationship with each other on LinkedIn (we will assume that relationships are symmetric). We are looking for a subset $S \subseteq V$ of vertices so that every vertex $s \in S$ has edges to at least 20 other vertices in S . And we want to make the set S as large as possible, subject to these constraints.

Design an efficient algorithm to find the set of super-schmoozers (the largest set S that is consistent with these constraints), given the graph G .

Hint: There are some vertices you can rule out immediately as not super-schmoozers.

2. (20 pts) A funky kind of coloring

Let $G = (V, E)$ be an undirected graph where every vertex has degree ≤ 51 . Let's find a way of coloring each vertex blue or gold, so that no vertex has more than 25 neighbors of its own color.

Consider the following algorithm, where we call a vertex "bad" if it has more than 25 neighbors of its own color:

1. Color each vertex arbitrarily.
2. Let $B := \{v \in V : v \text{ is bad}\}$.
3. While $B \neq \emptyset$:
4. Pick any bad vertex $v \in B$.
5. Reverse the color of v .
6. Update B to reflect this change, so that it again holds the set of bad vertices.

Notice that if this algorithm terminates, it is guaranteed to find a coloring with the desired property.

- (a) Prove that this algorithm terminates in a finite number of steps. I suggest that you define a *potential function* that associates a non-negative integer (the potential) to each possible way of coloring the graph, in such a way that each iteration of the while-loop is guaranteed to strictly reduce the potential.
- (b) Prove that the algorithm terminates after at most $|E|$ iterations of the loop.

Hint: You should figure out the largest value the potential could take on.

Optional: Think about how to implement the algorithm so that its total running time is $O(|V| + |E|)$ — this won't be graded.

3. (20 pts) Job Scheduling

You are given a set of n jobs, each runs in unit time. Job i has an integer-valued deadline time $d_i \geq 0$ and a real-valued penalty $p_i \geq 0$. Jobs may be scheduled to start at any non-negative integer time (0, 1, 2, etc), and only one job may run at a time. If job i completes at or before time d_i , then it incurs no penalty; otherwise, it is late and incurs penalty p_i . The goal is to schedule all jobs so as to minimize the total penalty incurred.

For each of the following greedy algorithms, either prove that it is correct, or give a simple counterexample (with at most three jobs) to show that it fails.

- Among unscheduled jobs that can be scheduled on time, consider the one whose deadline is the earliest (breaking ties with the highest penalty), and schedule it at the earliest available time. Repeat.
- Among unscheduled jobs that can be scheduled on time, consider the one whose penalty is the highest (breaking ties with the earliest deadline), and schedule it at the earliest available time. Repeat.
- Among unscheduled jobs that can be scheduled on time, consider the one whose penalty is the highest (breaking ties arbitrarily), and schedule it at the latest available time before its deadline. Repeat.

4. (20 pts) Timesheets

Suppose we have N jobs labelled $1, \dots, N$. For each job, you have determined the bonus of completing the job, $V_i \geq 0$, a penalty per day that you accumulate for not doing the job, $P_i \geq 0$, and the days required for you to successfully complete the job $R_i > 0$.

Every day, we choose one unfinished job to work on. A job i has been finished if we have spent R_i days working on it. This doesn't necessarily mean you have to spend R_i contiguous sequence of days working on job i . We start on day 1, and we want to complete all our jobs and finish with maximum reward. If we finish job i at the end of day t , we will get reward $V_i - t \cdot P_i$. Note, this value can be negative if you choose to delay a job for too long.

Given this information, what is the optimal job scheduling policy to complete all of the jobs?

5. (20 pts) Weighted Set Cover

In class (and Chapter 5.4) we looked at a greedy algorithm to solve the *set cover* problem, and proved that if the optimal set cover has size k , then our greedy algorithm will find a set cover of size at most $k \log_e n$.

Here is a generalization of the set cover problem.

- Input:** A set of elements B of size n ; sets $S_1, \dots, S_m \subseteq B$; positive weights w_1, \dots, w_m .
- Output:** A selection of the sets S_i whose union is B .
- Cost:** The sum of the weights w_i for the sets that were picked.

Design an algorithm to find the set cover with approximately the smallest cost. Prove that if there is a solution with cost k , then your algorithm will find a solution with cost $O(k \log n)$.